

Gaurav Gurjar

Senior Data Engineer

UAE Golden Visa holder

Based in UAE | English fluent ✉ ggurjar333@gmail.com in linkedin.com/in/ggurjarsoci GitHub @dataengineergaurav

PROFESSIONAL EXPERIENCE

6overn.ai (via Archetypal AI) | Technical Advisor, AI Data Platform 2026 → present
On-device AI browser agent that grounds answers to a verified knowledge base with no data leaving the device, delivered with Archetypal AI's NSA-grade GOVERN platform.

- Designed governed AI data pipelines that convert policies and knowledge sources into structured, searchable artifacts for model grounding.
- Built policy-aware LLM routing, KB-only execution paths, provenance controls, and governance telemetry for a multi-provider browser agent.
- Developed automated testing, live-browser QA, and release workflows improving reliability across AI-provider integrations.

Stack: policy-pipelines, LLM-routing, grounding, provenance, governance-telemetry, JavaScript, TypeScript, Python, pytest

ISHIR | Senior Data Engineer 2024-04 → 2026-01
Global outsourced IT-services and AI-native product-engineering firm founded 1999 in Dallas, with 1,400+ projects delivered across the US, India, and LATAM.

Client: SageSure
Largest independent US residential-property managing general underwriter, ~\$3.2B premium across 16 coastal states, 1,300+ employees.

The insurance data function built the dbt and Redshift platform behind the client's regulatory filings and business reporting.

- Owned the delivery end to end, including the reusable ETL packages and database modules underneath.
SageSure Insurance Data Platform
Fact and dimension models, compliance datasets, and analytical marts on Amazon Redshift using dbt and SQL, with data-quality tests and incremental snapshots.

Stack: dbt, Redshift, SQL, AWS, pytest

Petfolk USA | Consultant, Data Engineering (concurrent with ISHIR) 2025-10 → 2026-01
Vet-led veterinary-care company founded 2019 in Charlotte, \$40M Series B, combining in-person and virtual care through a membership model.

- Authored a modular inventory-adjustment ETL package covering 46 clinic locations using Airflow, boto3, Pydantic, and API integrations.
- Validated location coverage and mappings, handled data exceptions, and automated reproducible weekly releases.
Modular inventory-adjustment ETL with Airflow, Pydantic, and API integrations; validated mappings and weekly releases.

Stack: Airflow, FastAPI, AWS, Pydantic, boto3, API-integrations

CannaSpyGlass | Data Engineer 2021-04 → 2024-03
B2B regulated market-data and analytics platform founded 2021, serving retailers, multi-state operators, and investors including Leaflink, Harvard, and Prism Bank.

- Owned cloud ingestion and transformation for 300+ production pipelines spanning regulated-industry businesses, geospatial data, APIs, and PDF sources.
- Produced analytics-ready datasets with Python, SQL, AWS Glue, PySpark, GCP, and FastAPI services, including essential-services location data across the US.
Multi-source market-data platform consolidating wholesalers, transporters, producers, retailers, Google Maps, and PDF datasets across the United States.

Stack: AWS-Glue, PySpark, GCP, FastAPI, Python, PDF-parsing, geospatial

AI Squared Inc. | Data Scientist 2020-05 → 2021-03
Washington, DC enterprise-AI platform company, ~\$19.8M raised, whose Sparx and Unifi products embed AI/ML inside existing business applications.

- Built the statistical, data-science, and backend components for a public COVID-19 mortality-risk application, including the data pipelines and AWS Lambda/RDS-backed APIs it ran on, used by 2M+ people in the US.

COVID-19 Mortality Risk Tool
Statistical and data-science components in R/Python (wrapping the iCARE R package) with AWS Lambda/RDS-backed APIs for a public risk application.

Stack: R, Python, AWS-Lambda, RDS, iCARE

Casepoint LLC | Assistant Project Manager / Data & eDiscovery 2018-08 → 2020-04
AI-powered eDiscovery, FOIA, and compliance platform founded 2008 in Washington, DC, with DOD IL6 and FedRAMP authorizations, ~890 employees.

Handled data ingestion and processing for eDiscovery projects across law firms, enterprise organizations, and the U.S. Securities and Exchange Commission, under strict legal, regulatory, privacy, and evidentiary requirements where every ingested dataset had to be defensible.

- Managed ingestion of large litigation and investigation datasets using Casepoint's native ingestion tools and SQL, running validation, metadata verification, and processing checks so each dataset held up downstream and in discovery.
- Used SQL to investigate data issues, validate ingestion results, and reconcile records across multiple sources and file types, resolving ingestion exceptions, metadata problems, duplicates, and processing discrepancies.
- Supported \$1M+ eDiscovery projects while coordinating an 8-member delivery team, working directly with stakeholders across law firms, enterprise organizations, and the U.S. SEC.

Stack: SQL, eDiscovery, Casepoint-ingestion, metadata, data-validation

SUMMARY

Senior Data Engineer with 7+ years across data engineering, data science, and regulated eDiscovery delivery. Career arc from eDiscovery ingestion at a FedRAMP-authorized legal platform, to public-health data science used by 2M+ people, to 5 years of governed cloud data-platform work in insurance and regulated markets, including 300+ production pipelines built with AWS (Glue, Redshift, Lambda, S3), Python, SQL, dbt, PySpark, and Airflow. Strong in dimensional modeling, data quality, compliance datasets, and AI data governance.

INDUSTRIES

- AI Governance
- Insurance
- Public Health
- Veterinary Healthcare
- Regulated Market Data
- LegalTech / Regulated eDiscovery

SKILLS

Data Engineering: ETL/ELT, dimensional modeling, data quality, dbt, Airflow, Spark/PySpark, orchestration, data lineage

AI Data Systems: LLM routing, RAG, grounding/retrieval, policy pipelines, PII governance, telemetry

AWS & Cloud: S3, Glue, Lambda, Redshift, RDS, BigQuery

Languages & App Stack: Python, JavaScript, TypeScript, SQL, FastAPI, Pydantic, pytest, Docker, PostgreSQL, Pandas

EDUCATION

Bachelor of Computer Applications (BCA)

Dr. Babasaheb Ambedkar Open University

CERTIFICATIONS

IBM Data Engineer

IBM · 2025-01 · ID 61R9EX54ZNVG

IBM Data Warehouse Engineer

IBM · 2025-01 · ID 04EQFZ6OU7U

dbt Fundamentals

dbt Labs · 2025-02 · ID 134685079

Data Engineering Professional Certificate (V2)

Coursera · 2025-01 · ID H1LOTACMSSLG

Dagster & dbt

Dagster Labs · 2024-11 · ID scmduvikqm

UAE AI Camp Certificate 2024

Artificial Intelligence Office (UAE) · 2024-07

Certified Prompt Engineering

Promptmaster · 2024-05