

Gaurav Gurjar

Senior Data Engineer

UAE Golden Visa holder

Based in UAE · ggurjar333@gmail.com · linkedin.com/in/ggurjarsocl · GitHub @dataengineergaurav

SUMMARY

Senior Data Engineer with 7+ years across data engineering, data science, and regulated eDiscovery delivery. Career arc from eDiscovery ingestion at a FedRAMP-authorized legal platform, to public-health data science used by 2M+ people, to 5 years of governed cloud data-platform work in insurance and regulated markets, including 300+ production pipelines built with AWS (Glue, Redshift, Lambda, S3), Python, SQL, dbt, PySpark, and Airflow. Strong in dimensional modeling, data quality, compliance datasets, and AI data governance.

PROFESSIONAL EXPERIENCE

Senior Data Engineer | ISHIR (client: US residential property insurer)

2024-04 → 2026-01

IT Services / Data & AI · employee

Global outsourced IT-services and AI-native product-engineering firm founded 1999 in Dallas, with 1,400+ projects delivered across the US, India, and LATAM.

Client: US residential property insurer

Largest independent US residential-property managing general underwriter, ~\$3.2B premium across 16 coastal states, 1,300+ employees.

The insurance data function built the dbt and Redshift platform behind the client's regulatory filings and business reporting.

- Owned the delivery end to end, including the reusable ETL packages and database modules underneath.

Insurance Data Platform

- Developed fact and dimension tables, compliance datasets, and analytical data marts on Amazon Redshift using dbt and SQL.
- Implemented dbt data-quality tests and SQL integration tests; deployed daily and monthly incremental snapshots with scheduled Redshift automation.

dbt Redshift SQL AWS pytest

Consultant, Data Engineering | US veterinary-care operator (concurrent with ISHIR)

2025-10 → 2026-01

Veterinary Healthcare · consultant

Vet-led veterinary-care company founded 2019 in Charlotte, \$40M Series B, combining in-person and virtual care through a membership model.

- Authored a modular inventory-adjustment ETL package covering 46 clinic locations using Airflow, boto3, Pydantic, and API integrations.
- Validated location coverage and mappings, handled data exceptions, and automated reproducible weekly releases.

Airflow FastAPI AWS Pydantic boto3 API-integrations

Data Engineer | CannaSpyGlass

2021-04 → 2024-03

Regulated Market Data & Analytics · contractor

B2B regulated market-data and analytics platform founded 2021, serving retailers, multi-state operators, and investors including Leaflink, Harvard, and Prism Bank.

- Owned cloud ingestion and transformation for 300+ production pipelines spanning regulated-industry businesses, geospatial data, APIs, and PDF sources.
- Produced analytics-ready datasets with Python, SQL, AWS Glue, PySpark, GCP, and FastAPI services, including essential-services location data across the US.

AWS-Glue PySpark GCP FastAPI Python PDF-parsing geospatial

Data Scientist | AI Squared Inc.

2020-05 → 2021-03

Enterprise AI Platform · contractor

Washington, DC enterprise-AI platform company, ~\$19.8M raised, whose Sparx and Unifi products embed AI/ML inside existing business applications.

- Built the statistical, data-science, and backend components for a public COVID-19 mortality-risk application, including the data pipelines and AWS Lambda/RDS-backed APIs it ran on, used by 2M+ people in the US.

COVID-19 Mortality Risk Tool

- Developed statistical and data-science components in R and Python, wrapping iCARE R package functions in Python.
- Built AWS Lambda and RDS-backed data services for COVID19RiskTools.com.

R Python AWS-Lambda RDS iCARE

Assistant Project Manager / Data & eDiscovery | Casepoint LLC

2018-08 → 2020-04

LegalTech / Regulated eDiscovery · employee

AI-powered eDiscovery, FOIA, and compliance platform founded 2008 in Washington, DC, with DOD IL6 and FedRAMP authorizations, ~890 employees.

Handled data ingestion and processing for eDiscovery projects across law firms, enterprise organizations, and the U.S. Securities and Exchange Commission, under strict legal, regulatory, privacy, and evidentiary requirements where every ingested dataset had to be defensible.

- Managed ingestion of large litigation and investigation datasets using Casepoint's native ingestion tools and SQL, running validation, metadata verification, and processing checks so each dataset held up downstream and in discovery.
- Used SQL to investigate data issues, validate ingestion results, and reconcile records across multiple sources and file types, resolving ingestion exceptions, metadata problems, duplicates, and processing discrepancies.
- Supported \$1M+ eDiscovery projects while coordinating an 8-member delivery team, working directly with stakeholders across law firms, enterprise organizations, and the U.S. SEC.

SQL eDiscovery Casepoint-ingestion metadata data-validation

INDUSTRIES

Insurance Public Health Veterinary Healthcare Regulated Market Data LegalTech / Regulated eDiscovery

SKILLS

Data Engineering: ETL/ELT, dimensional modeling, data quality, dbt, Airflow, Spark/PySpark, orchestration, data lineage
AWS & Cloud: S3, Glue, Lambda, Redshift, RDS, BigQuery
Languages & App Stack: Python, JavaScript, TypeScript, SQL, FastAPI, Pydantic, pytest, Docker, PostgreSQL, Pandas

EDUCATION

- **Master of Science, Data Science**, Institute of Advanced Research, Gandhinagar, Dropout
- **Bachelor of Computer Applications (BCA)**, Dr. Babasaheb Ambedkar Open University

CERTIFICATIONS

- **IBM Data Engineer**, IBM, 2025-01 (ID 61R9EX54ZNVG)
- **IBM Data Warehouse Engineer**, IBM, 2025-01 (ID 04EJQFZ6OU7U)
- **dbt Fundamentals**, dbt Labs, 2025-02 (ID 134685079)
- **Data Engineering Professional Certificate (V2)**, Coursera, 2025-01 (ID H1LOTACMSSLC)
- **Dagster & dbt**, Dagster Labs, 2024-11 (ID scmdubikqm)
- **UAE AI Camp Certificate 2024**, Artificial Intelligence Office (UAE), 2024-07
- **Certified Prompt Engineering**, Promptmaster, 2024-05

CLIENT WORK

Insurance Data Platform US residential property insurer · Insurance (Catastrophe-exposed Property)
Fact and dimension models, compliance datasets, and analytical marts on Amazon Redshift using dbt and SQL, with data-quality tests and incremental snapshots.

Redshift dbt SQL

CannaSpyGlass Regulated Market Data Platform CannaSpyGlass · Regulated Market Data & Analytics
Multi-source market-data platform consolidating wholesalers, transporters, producers, retailers, Google Maps, and PDF datasets across the United States.

AWS-Glue PySpark GCP FastAPI

COVID-19 Mortality Risk Tool COVID19RiskTools.com (Public Health) · Public Health
Statistical and data-science components in R/Python (wrapping the iCARE R package) with AWS Lambda/RDS-backed APIs for a public risk application.

R Python AWS-Lambda RDS

eDiscovery Ingestion & Delivery Casepoint LLC · LegalTech / Regulated eDiscovery
Ingestion and processing of litigation and investigation datasets for law firms, enterprise organizations, and the U.S. Securities and Exchange Commission, using Casepoint's native ingestion tooling and SQL for validation, metadata verification, and record reconciliation. Work ran under legal, regulatory, privacy, and evidentiary requirements, so every ingested dataset had to be defensible downstream and in discovery.
Stoel Rives · Robbins Russell · U.S. SEC · Turner Construction USA · Beveridge & Diamond · Briggs & Morgan

SQL eDiscovery Casepoint-ingestion metadata data-validation

Veterinary Inventory ETL US veterinary-care operator · Veterinary Healthcare
Modular inventory-adjustment ETL with Airflow, Pydantic, and API integrations; validated mappings and weekly releases.

Airflow FastAPI Pydantic AWS